

Algorithms and Data Structures with Applications in Machine Learning

Graph Representation Learning



January 1, 2025



Graph Terminology and Representation

Graph Representation Learning: DeepWalk and Node2Vec

Graph Neural Networks



Graph Terminology and Representation

Graph Representation Learning: DeepWalk and Node2Vec

Graph Neural Networks



Definition

A graph is defined as:

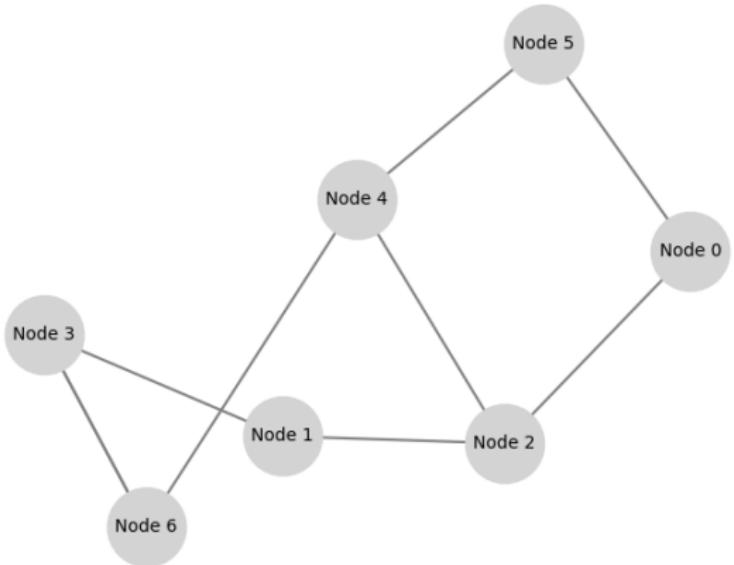
$$G = (V, E, u)$$

- ▶ **Nodes (Vertices):** The set V represents the nodes in the graph.
- ▶ **Edges:** The set $E \subseteq V \times V$ represents the connections (relationships) between the nodes.
- ▶ **Features:** Each node can have a feature vector $u(v)$ representing its attributes.
- ▶ **Labels:** Nodes (or edges) can also have labels, which are used for tasks like classification.

Example Graph



Example: The graph below has 7 connected nodes ($V = \{0, 1, 2, 3, 4, 5, 6\}$) and their edges (E).

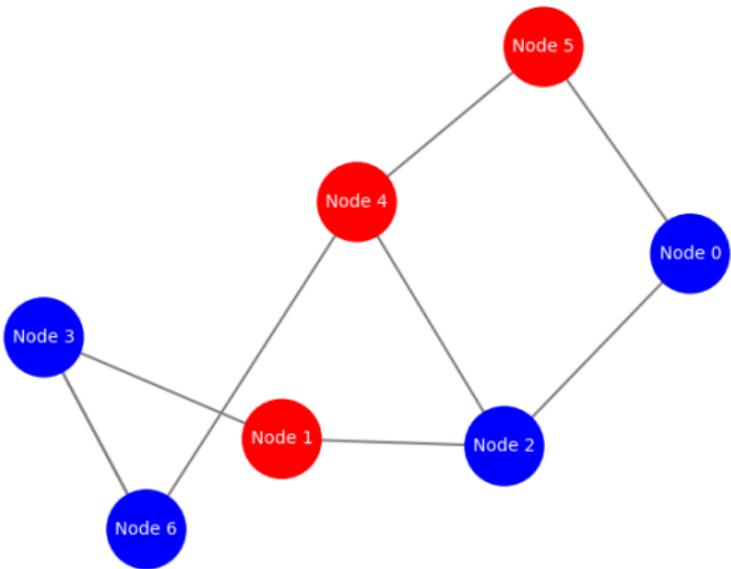


Example Graph: Node Labels



Example: Nodes in a graph can be associated with labels.


Blue nodes: Label 0 **Red nodes:** Label 1



Example Graph: Node Features



Example: Each node in the graph can have associated features. In this case: Each node has a feature vector of dimension 3.



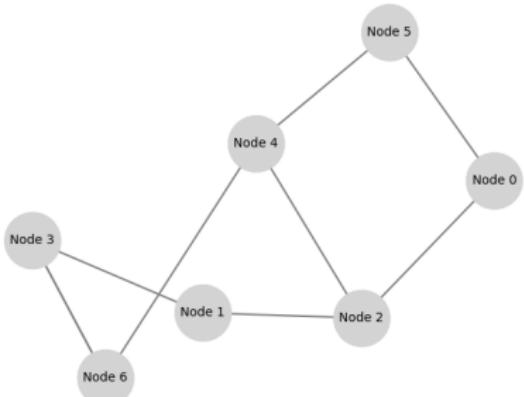
Adjacency Matrix

Definition

The adjacency matrix A of a graph $G = (V, E)$ is a matrix of size $|V| \times |V|$, where:

- ▶ $A[i][j] = 1$ if there is an edge between node i and node j .
- ▶ $A[i][j] = 0$ if there is no edge between node i and node j .

Example: A graph and its corresponding adjacency matrix:



Adjacency Matrix:

$$A = \begin{bmatrix} 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 1 & 1 & 0 & 0 & 0 \\ 1 & 1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 & 1 & 0 & 0 \end{bmatrix}$$


Weighted Adjacency Matrix

Definition

The adjacency matrix A can be extended to a weighted matrix W , where:

- ▶ $W[i][j]$ represents the weight of the edge between node i and node j .

Example: A graph and its a weighted adjacency matrix:



Weighted Matrix:

$$W = \begin{bmatrix} 0 & 0 & 2.1 & 0 & 0 & 1.5 & 0 \\ 0 & 0 & 3.4 & 1.8 & 0 & 0 & 0 \\ 2.1 & 3.4 & 0 & 0 & 2.5 & 0 & 0 \\ 0 & 1.8 & 0 & 0 & 0 & 0 & 4.2 \\ 0 & 0 & 2.5 & 0 & 0 & 1.2 & 3.0 \\ 1.5 & 0 & 0 & 0 & 1.2 & 0 & 0 \\ 0 & 0 & 0 & 4.2 & 3.0 & 0 & 0 \end{bmatrix}$$



Applications: Machine Learning on graphs enables a variety of tasks, including:

- ▶ **Node Prediction:** Predict properties or labels of nodes in a graph (e.g., user classification in social networks).
- ▶ **Link Prediction:** Predict the existence or strength of a connection between two nodes (e.g., recommendation systems).
- ▶ **Graph Classification:** Assign labels to entire graphs (e.g., chemical compound classification).
- ▶ **Clustering:** Group nodes into communities or clusters based on their properties or structure.

Objective: Node Classification




Objective: The objective of this course is two-fold:

- 1. Learning a D -dimensional representation:**

Create embedding vectors for nodes that capture the structure of the graph.

- 2. Node Classification:**

Use the learned embeddings to predict the labels of the nodes.





Graph Terminology and Representation

Graph Representation Learning: DeepWalk and Node2Vec

Graph Neural Networks



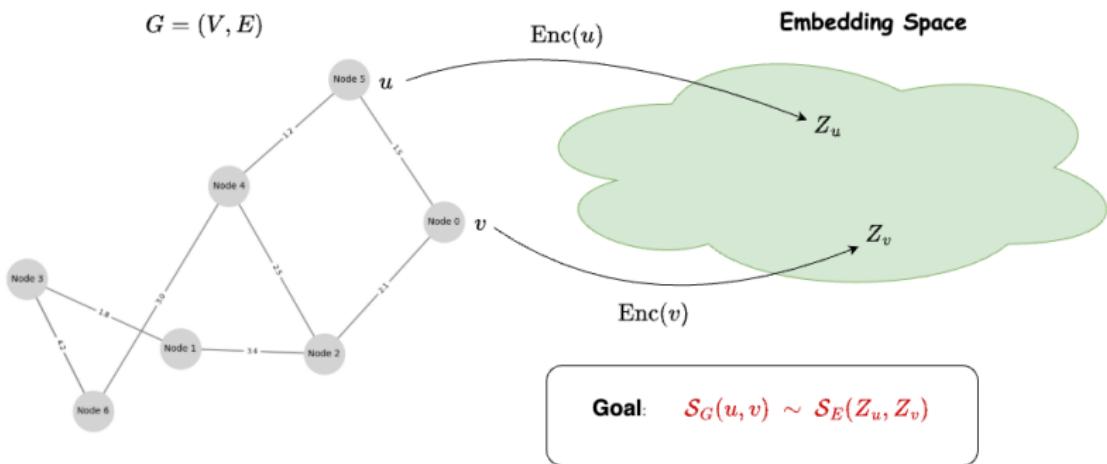
Objective: We aim to learn a mapping:

$$f : V \rightarrow \mathbb{R}^D$$

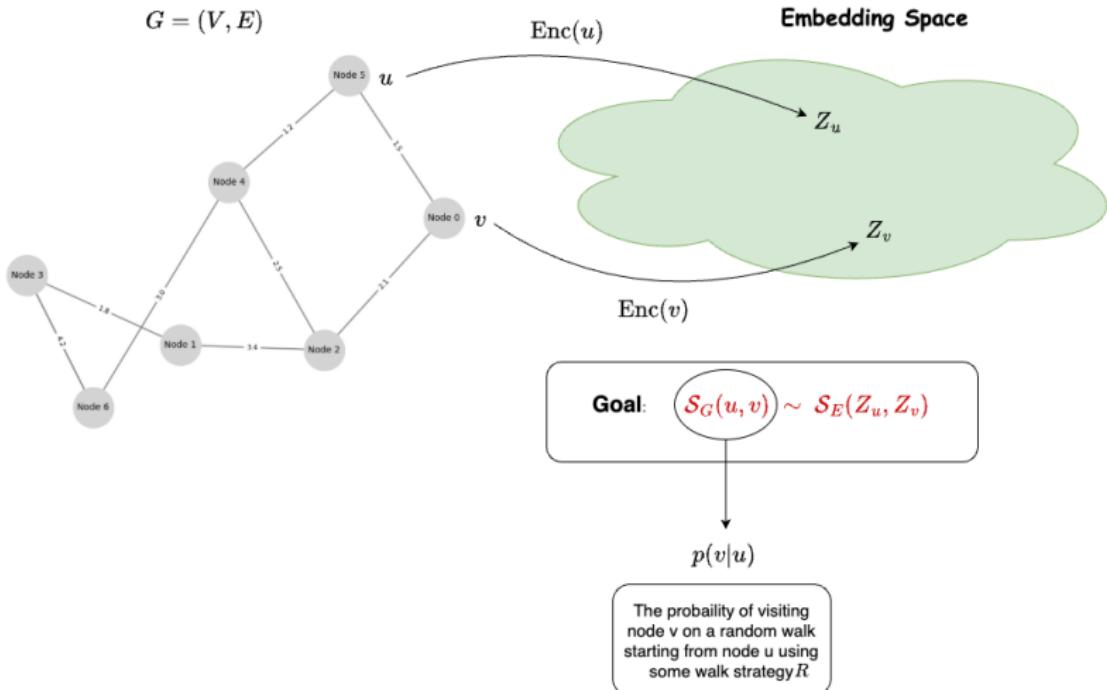
where each node $u \in V$ is mapped to a D -dimensional vector $\mathbf{z}_u \in \mathbb{R}^D$.

- ▶ In this section, we focus on leveraging the graph's **structure** to generate embedding vectors for nodes.
- ▶ The embeddings can be used for downstream tasks, such as node classification or link prediction.
- ▶ **No use of feature vectors:** We only use the graph topology (connections between nodes) to derive the embeddings.

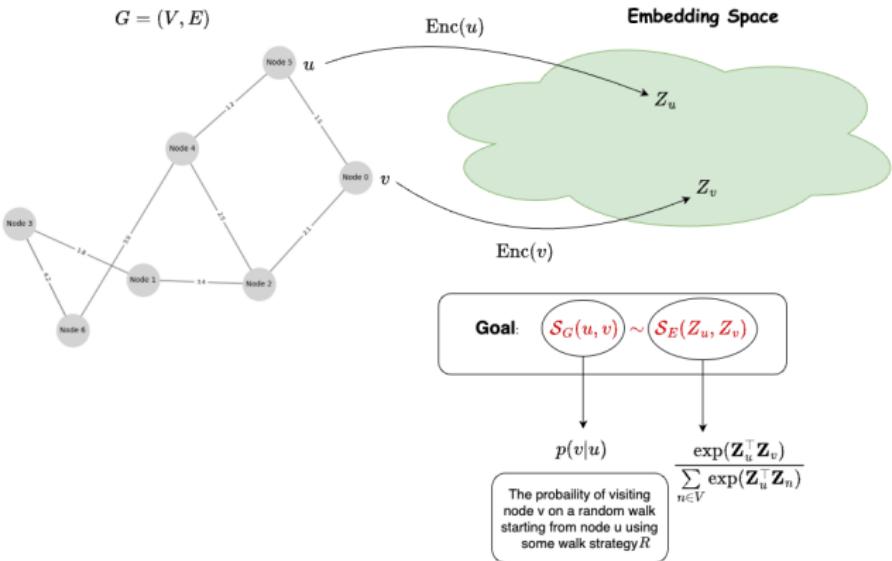
Graph Structure-Based Embeddings: Objective



Graph Structure-Based Embeddings: Objective




Graph Structure-Based Embeddings: Objective



Random Walks:

- ▶ A random walk is a sequence of steps through the graph, starting from a given node u , where each step randomly selects a neighboring node.
- ▶ The nodes visited during these walks represent the local neighborhood structure around u , denoted $\mathcal{N}_R(u)$
- ▶ Here is an example of a random walk from node u to node v .





Algorithm Fixed-Length Random Walks

Require: Graph $G = (V, E)$, starting node u , walk length L , number of walks N

Ensure: $\mathcal{N}_R(u)$ Multiset of nodes visited during random walks starting from u

```
1: Initialize an empty multiset of neighbors: neighbors ← []
2: for  $n = 1$  to  $N$  do                                ▷ Perform  $N$  random walks
3:     Initialize current_node ←  $u$ 
4:     for  $l = 1$  to  $L$  do                      ▷ Walk for  $L$  steps
5:         Sample a random neighbor  $v \in \text{Neighbors}(\text{current\_node})$ 
6:         neighbors.append( $v$ )
7:         current_node ←  $v$ 
8:     end for
9: end for
10: return neighbors
```

Introducing Node2Vec: Biased Random Walks



- ▶ The Node2Vec algorithm modifies traditional random walks by introducing **biases** that control how the walk explores the graph.
- ▶ This bias allows us to interpolate between two extremes:
 1. **Local Behavior:** Tendency to return to previously visited nodes, capturing local neighborhood structures. This is controlled by the **return hyperparameter p** .
 2. **Global Behavior:** Tendency to explore new, distant nodes, capturing the global structure of the graph. This is controlled by the **in-out hyperparameter q** .
- ▶ By adjusting p and q , Node2Vec generates embeddings that can reflect different graph traversal strategies.
- ▶ This flexibility makes Node2Vec suitable for capturing diverse graph structures. (See Programming Session 6).

Introducing Node2Vec: Biased Random Walks




- ▶ When the walk moves from node u to w , the neighbors of w are categorized based on their distance to u .
- ▶ We define the following **unnormalized probabilities**:
 1. Nodes closer to u than w receive an unnormalized probability of $\frac{1}{p}$.
 2. Nodes farther from u than w receive an unnormalized probability of $\frac{1}{q}$.
 3. Nodes at the same distance as w from u receive an unnormalized probability of 1.
- ▶ These unnormalized probabilities are normalized to form a valid probability distribution, which guides the biased random walk.

Introducing Node2Vec: Biased Random Walks



Here is an example of assigning the unnormalized probabilities:

- ▶ Starting at node u , the walk reaches node w .
- ▶ The probabilities assigned to w 's neighbors depend on their distance to u , as described in the previous slide.





Algorithm Biased Random Walks

Require: Graph $G = (V, E)$, starting node u , walk length L , number of walks N , return parameter p , in-out parameter q

Ensure: $\mathcal{N}_R(u)$: Multiset of nodes visited during biased random walks starting from u

```
1: Initialize an empty multiset of neighbors: neighbors ← []
2: for  $n = 1$  to  $N$  do           ▷ Perform  $N$  biased random walks
3:   Initialize current_node ←  $u$  and prev_node ← None
4:   for  $l = 1$  to  $L$  do           ▷ Walk for  $L$  steps
5:     Compute probabilities using prev_node and current_node
6:     Sample the next node  $v$  based on the these probabilities
7:     neighbors.append( $v$ )
8:     Update prev_node and current_node
9:   end for
10: end for
11: return neighbors
```



Defining the Loss Function:

- ▶ Now that we know how to define $\mathcal{N}_R(u)$, we can derive the loss function to train the embeddings.
- ▶ The objective is to minimize the following loss function:

$$\mathcal{L}(\theta) = - \sum_{u \in V} \sum_{v \in \mathcal{N}_R(u)} \log \left(\frac{\exp(\mathbf{Z}_u^\top \mathbf{Z}_v)}{\sum_{n \in V} \exp(\mathbf{Z}_u^\top \mathbf{Z}_n)} \right)$$

Where:

- ▶ $\mathbf{Z}_i \in \mathbb{R}^D$ is the embedding vectors for nodes $i \in V$.
- ▶ $\theta = \{\mathbf{Z}_i \mid i \in V\}$ represents all the embedding parameters to be learned.

Outline



Graph Terminology and Representation

Graph Representation Learning: DeepWalk and Node2Vec

Graph Neural Networks



Node2vec recap:

- ▶ Node2Vec generates embeddings by combining graph topology and biased random walks.
- ▶ Focuses solely on the graph structure, without leveraging node-specific feature vectors.

Paradigm Shift:

- ▶ Our new objective is to incorporate both **graph structure** and **node features** into the embeddings.
- ▶ Instead of manually defining the impact of neighbors (e.g., via p and q), we aim for the model to **learn** the importance of different neighbors.



Notations:

- ▶ $\mathbf{h}_v^{(k)}$: Learned embedding of node v at iteration k .
- ▶ $\mathcal{N}(v)$: Set of neighbors of node v .

At each iteration, embeddings are refined by aggregating information from the local neighborhood and updating the node's representation.

Steps for One Iteration (k):

1. **Aggregation:** Gather information from neighbors of node v :

$$\mathbf{a}_v^{(k)} = f_{\text{aggregate}} \left(\{\mathbf{h}_u^{(k-1)} \mid u \in \mathcal{N}(v)\} \right)$$

2. **Update:** Combine aggregated information and the previous embedding to compute the new embedding:

$$\mathbf{h}_v^{(k)} = f_{\text{update}}(\mathbf{a}_v^{(k)}, \mathbf{h}_v^{(k-1)})$$



Algorithm Message Passing Framework

Require: Graph $G = (V, E)$, node features $\{\mathbf{x}_v \mid v \in V\}$, number of iterations K , $f_{\text{aggregate}}$, f_{update}

Ensure: Final node embeddings $\{\mathbf{h}_v^{(K)} \mid v \in V\}$

1: Initialize embeddings: $\mathbf{h}_v^{(0)} \leftarrow \mathbf{x}_v$ for all $v \in V$

2: **for** $k = 1$ to K **do**

3: **for** each node $v \in V$ **do**

$$\mathbf{a}_v^{(k)} \leftarrow f_{\text{aggregate}} \left(\{\mathbf{h}_u^{(k-1)} \mid u \in \mathcal{N}(v)\} \right)$$


$$\mathbf{h}_v^{(k)} \leftarrow f_{\text{update}}(\mathbf{a}_v^{(k)}, \mathbf{h}_v^{(k-1)})$$

4: **end for**

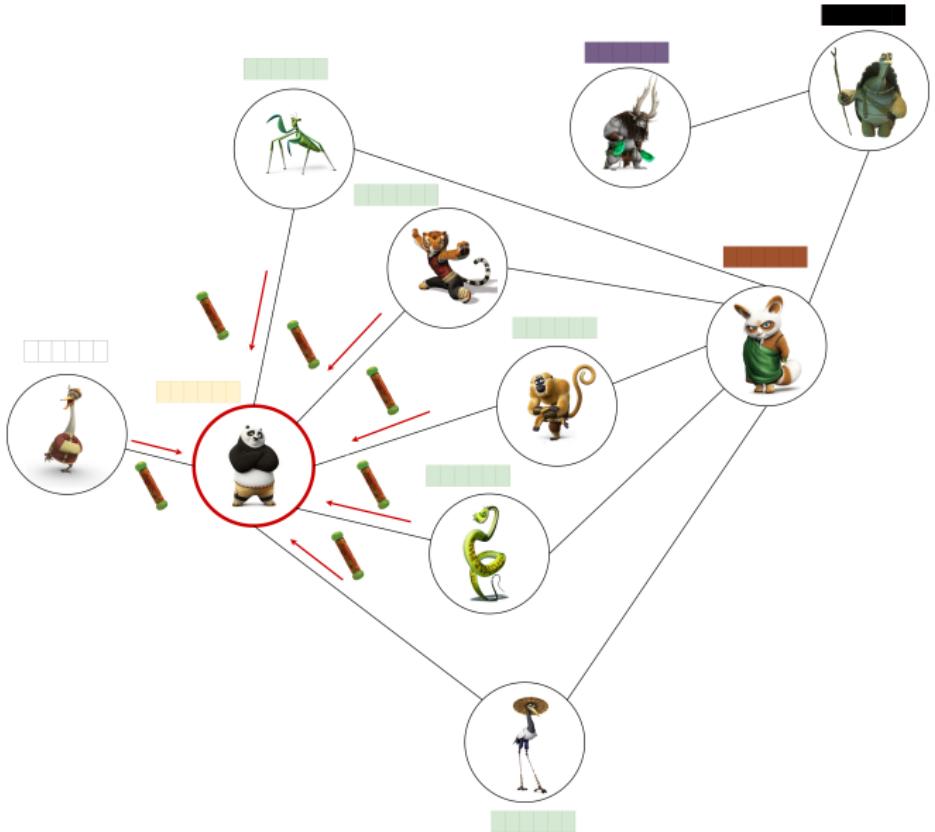
5: **end for**

6: **return** $\{\mathbf{h}_v^{(K)} \mid v \in V\}$


An Example: Graph Initialization with Feature Vectors



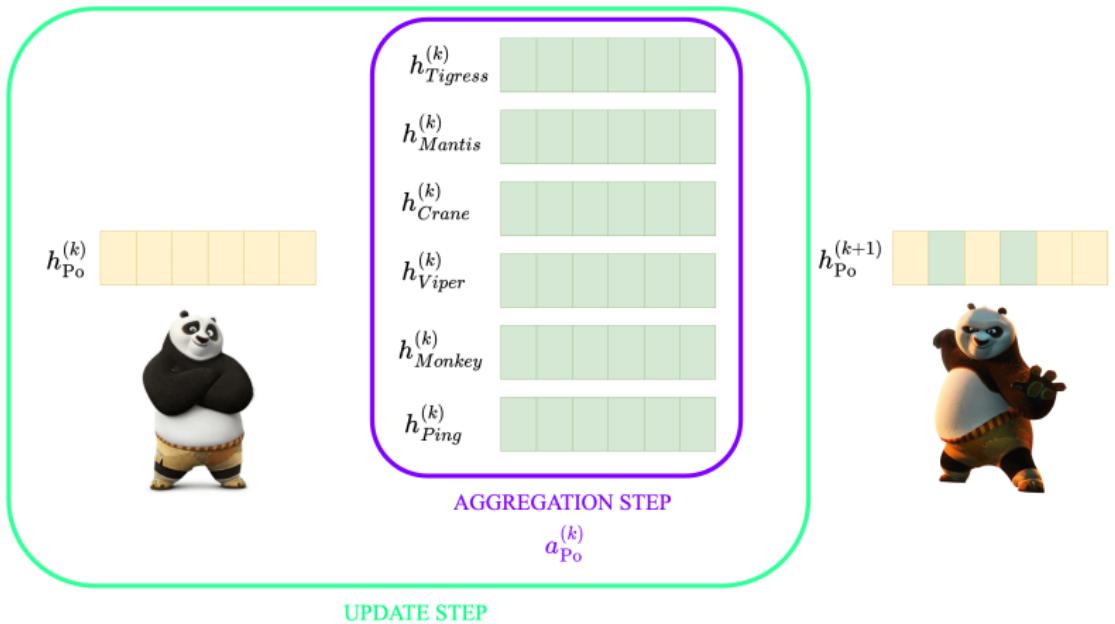
An Example: Aggregation Step





An Example: Update Step



An Example: Recap of Both Steps



An Example: Final Embedding Vectors





Aggregation Function ($f_{\text{aggregate}}$):

$$\mathbf{a}_v^{(k)} = \frac{1}{\deg(v)} \sum_{u \in \mathcal{N}(v)} \mathbf{h}_u^{(k-1)}$$

Update Function (f_{update}):

$$\mathbf{h}_v^{(k)} = \sigma \left(W^{(k)} \cdot \left[\mathbf{h}_v^{(k-1)} \parallel \mathbf{a}_v^{(k)} \right] \right)$$

Description:

- ▶ Aggregates the mean of the neighbors' embeddings.
- ▶ Updates the embedding with a learned linear transformation using weights $W^{(k)}$ and a non-linear activation σ (e.g., ReLU).



Aggregation Function ($f_{\text{aggregate}}$):

$$\mathbf{a}_v^{(k)} = \sum_{u \in \mathcal{N}(v) \cup \{v\}} \frac{\mathbf{h}_u^{(k-1)}}{\sqrt{\deg(v) \cdot \deg(u)}}$$

Update Function (f_{update}):

$$\mathbf{h}_v^{(k)} = \sigma \left(W^{(k)} \cdot \mathbf{a}_v^{(k)} \right)$$

Description:

- ▶ **Aggregation:** Aggregates information from neighbors and the node itself, normalized by the degree of both nodes.
- ▶ **Update:** Applies a linear transformation using $W^{(k)}$, followed by a non-linear activation σ (e.g., ReLU).



Aggregation Function ($f_{\text{aggregate}}$):

$$\mathbf{a}_v^{(k)} = \sum_{u \in \mathcal{N}(v) \cup \{v\}} \alpha_{vu} \mathbf{h}_u^{(k-1)}$$

Where:

$$\alpha_{vu} = \frac{\exp \left(\text{LeakyReLU} \left(\mathbf{a}^\top \left[\mathbf{h}_v^{(k-1)} \| \mathbf{h}_u^{(k-1)} \right] \right) \right)}{\sum_{w \in \mathcal{N}(v) \cup \{v\}} \exp \left(\text{LeakyReLU} \left(\mathbf{a}^\top \left[\mathbf{h}_v^{(k-1)} \| \mathbf{h}_w^{(k-1)} \right] \right) \right)}$$

Description:

- ▶ **Aggregation:** Computes a weighted sum of neighbor embeddings using attention coefficients α_{vu} .
- ▶ **Attention Coefficients α_{vu} :** Learn to assign importance to each neighbor dynamically.



Update Function (f_{update}):

$$\mathbf{h}_v^{(k)} = \left\|_{k=1}^K \sigma \left(W_k^{(k)} \mathbf{a}_v^{(k)} \right) \right\|$$

Description:

- ▶ **Multi-Head Attention:** Combines results from K independent attention heads by concatenation ($\|$).
- ▶ **Non-Linearity:** Applies a learned linear transformation $W_k^{(k)}$ followed by a non-linear activation σ (e.g., ReLU).
- ▶ GATs allow each node to focus on the most relevant neighbors dynamically, enabling better representation learning for tasks such as node classification or graph-level predictions.



Objective: Train node embeddings $\mathbf{h}_v^{(K)}$ by leveraging the graph structure, without requiring labels.

The Loss Function:

$$\mathcal{L}(\theta) = - \sum_{u \in V} \sum_{v \in \mathcal{N}_R(u)} \log \left(\frac{\exp(\mathbf{h}_u^{(K)\top} \mathbf{h}_v^{(K)})}{\sum_{n \in V} \exp(\mathbf{h}_u^{(K)\top} \mathbf{h}_n^{(K)})} \right)$$

Where:

- ▶ $\mathbf{h}_u^{(K)}$: Final embedding of node u after K message-passing layers.
- ▶ $\mathcal{N}_R(u)$: Neighborhood of u defined using some random walk strategy.
- ▶ We usually approximate the denominator using **negative sampling**.



Objective: Predict the label of each node $v \in \mathcal{V}_{\text{train}}$ using the GNN-generated embeddings $\mathbf{h}_v^{(K)}$.

The Loss Function (Cross-Entropy):

$$\mathcal{L} = - \sum_{v \in \mathcal{V}_{\text{train}}} \sum_{c=1}^C y_v^c \log \hat{y}_v^c$$

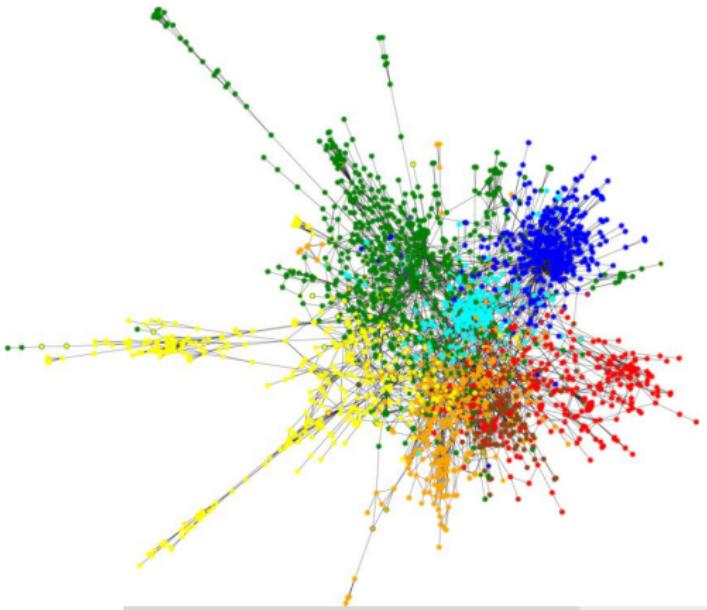
Where:

- ▶ y_v^c : Ground-truth label (one-hot encoded) for node v .
- ▶ $\hat{y}_v^c = \text{softmax}\left(W_{\text{out}} \mathbf{h}_v^{(K)}\right)$: Predicted probability of class c , computed from the node embedding.

Programming Session: Node Classification



- ▶ During the programming session, we will work on the **Cora dataset**.
- ▶ The objective will be to build and train a **Graph Neural Network (GNN)** for node classification.



Thank you for your attention